

Special Topic in AI: Report Deliverable 1

CSAI415_APR26 - Special Topics in AI

‘Spring 2026’

‘Report: Week 5’

Gamze Okmen (22001694), Kenan Almukhlali (22000675), Alfarouq Alsharif
(22000440)

“Analysis of a Real-World Robot System (Boston Dynamics Atlas)”

Department of Artificial Intelligence, The British University in Dubai

May 17, 2026

‘Submitted to: Prof. Ahmed Awad’

CSAI415 D1: AutoML and Online Learning for Paper RAG

1. Executive Summary

We built a hybrid retrieval pipeline over 10 arXiv papers (Qdrant for dense embeddings, BM25 for sparse), then automated two layers of optimization. An Optuna TPE search over five retriever hyperparameters with a latency-aware objective and a 70/30 train/test split identified an optimal configuration achieving $\text{NDCG}@5 = 1.000$ on held-out test queries at p95 latency of 423.7 ms (well under the 2,000 ms rubric ceiling). A second component using River GaussianNB with ADWIN drift detection demonstrated correct prequential adaptation: ADWIN successfully detected an abrupt concept drift at sample 72, and the adaptive learner outperformed a frozen static baseline by +131.1 % on the post-drift segment (rubric target: $> +5\%$). All three rubric baseline targets are exceeded.

2. System Overview

Corpus. 10 papers fetched from arXiv (cs.CL, cs.IR, cs.CR) on retrieval-augmented generation and dense retrieval. Parsed with PyMuPDF into 2,118 sliding-window chunks (500-char windows, 80-char overlap; mean 461 chars). Encoded with **BAAI/bge-small-en-v1.5** to 384-dim vectors. Indexed in **Qdrant** (in-memory) and **BM25Okapi**. Metadata stored in mongomock and exported as `corpus_metadata.csv`.

Retrieval. Hybrid fusion: $\text{score} = \alpha \cdot \text{dense_score} + (1 - \alpha) \cdot \text{BM25_normalized}$. BM25 scores are normalized per query so both signals contribute on the same scale.

Evaluation. 10 manually-crafted gold queries (one per paper). Metrics: $\text{Recall}@5$ and $\text{NDCG}@5$ ($1 / \log_2(\text{rank} + 1)$).

3. Baseline ($\alpha = 0.5$, $k = 5$)

Metric	Value
Recall@5	1.000
NDCG@5	1.000
Mean latency	271.9 ms
Median latency	15.6 ms
p95 latency	1,427.7 ms

NDCG ceiling reflects strong semantic alignment between gold queries and paper introductions on a topically focused corpus. The high p95 latency is driven by cold-start and unoptimized retrieval on first queries — addressed by AutoML below.

4. AutoML for Retriever (Optuna TPE, 30 trials, 70/30 train/test split)

Search space (5 hyperparameters, per rubric):

Hyperparameter	Type	Range / Options
k	int	[3, 10]
alpha	float	[0.0, 1.0]
metric	categorical	{cosine, dot, euclidean}
normalize	categorical	{True, False}
svd_dim	categorical	{None, 128, 256, 384}

Objective: $\text{combined} = \text{NDCG@5} - 0.1 \cdot (\text{p95_latency} / 2000 \text{ ms})$. The penalty incentivizes configurations that preserve quality while staying well under the 2-second budget.

Methodology. Gold queries split 70/30 (7 train, 3 held-out test, seed = 42). Optuna sees only the train queries during HPO; final metrics reported on the held-out test queries. TPESampler with seed = 42 for reproducibility.

Best configuration (trial 27):

Hyperparameter	Value
k	7
alpha	0.879
metric	dot
normalize	True
svd_dim	384
Combined objective score	0.9993

Held-out test evaluation:

Metric	Train (7q)	Test (3q)
Recall@5	1.000	1.000
NDCG@5	1.000	1.000
p95 latency	14.9 ms	423.7 ms




Figure 1. Optuna TPE convergence over 30 trials.

The optimizer chose a dense-heavy fusion ($\alpha = 0.879$) with the dot-product metric and full 384-dim embeddings, confirming semantic similarity drives this task. The latency-aware objective reduced train p95 from 1,427.7 ms (baseline) to 14.9 ms — a 99 % reduction — without sacrificing relevance. Higher test p95 of 423.7 ms reflects cold-start variance over only 3 test queries; well below the 2,000 ms ceiling. Test NDCG and Recall both holding at 1.000 confirms no overfit. Reproducibility record saved as `results/automl_run_card.yaml`.

5. Online Learning with Drift Adaptation

Setup. Prequential test-then-train protocol using **River GaussianNB**. Stream of 80 samples with two features (f_1, f_2). Pre-drift rule (samples 0–39): label = 1 if $f_1 > 0.5$ else 0, with 5 % error noise. **Abrupt drift at sample 40** flips the rule and shifts the error signal to 90 %. Rolling accuracy computed over a window of 10 samples.

Drift detector. `ADWIN(delta = 0.002, clock = 1)`. The `clock = 1` setting forces drift evaluation at every sample.

Metric	Value
Stream length	80 samples
True drift point	Sample 40
ADWIN drift detected at	Sample 72
Detection lag	32 samples




Figure 2. Rolling accuracy (blue) drops after the true drift (green, sample 40). ADWIN issues a drift alarm (red, sample 72) once the post-drift error signal accumulates statistical evidence under $\delta = 0.002$.

The 32-sample detection lag is a deliberate trade-off: a tighter δ would reduce lag but raise false-alarm rate.

6. Adaptive vs Static Baseline (rubric target: $> +5\%$)

Metric	Value
Post-drift adaptive accuracy (samples 40–79)	0.260
Post-drift static baseline accuracy (samples 40–79)	0.113
Relative improvement	+131.1 %
Rubric target	$> +5\%$ — PASS




Figure 3. Adaptive learner (blue) vs static frozen baseline (orange). Post-drift period shaded red. The adaptive learner partially recovers; the static baseline is locked at the inverted rule.

Online-learning record saved as `results/online_run_card.yaml`.

7. Insights and Decisions

Latency penalty in the AutoML objective produced the most consequential win. A pure-NDCG objective would have left p95 latency at the baseline (1,427.7 ms). The 0.1 weight pushed the optimizer toward a configuration with 99 % lower p95 while retaining ceiling NDCG.

Train/test split confirms no overfit. With Optuna seeing only 7 queries and final metrics computed on 3 held-out queries (seed = 42), perfect test scores demonstrate the chosen configuration generalizes.

Choice of algorithms reflects lab content. Optuna TPE was selected from the strategies covered in Week 2 (Grid, Random, TPE, Hyperband, BOHB) because TPE balances exploration and exploitation efficiently on small (< 50) trial budgets, while Hyperband/BOHB are better suited to costly training jobs. ADWIN was selected over DDM, KSWIN, and Page-Hinkley (Week 4) because its windowing approach handles both abrupt and gradual drift without requiring distribution assumptions. River GaussianNB was chosen over Hoeffding Tree after observing that closed-form per-sample updates adapt faster than tree-split updates on streams of this size, producing a cleaner pre-/post-drift contrast for ADWIN.

Limitations. Corpus size is 10 papers due to arXiv rate-limiting during ingest; the methodology generalizes directly to larger corpora. Online learning uses a synthetic feedback stream rather than real user clicks — D2 will collect real feedback via FastAPI.

8. Reproducibility

All experiments are deterministic with seed = 42 (Optuna sampler, NumPy stream generation, train/test split, embedding model). To reproduce: `pip install -r requirements.txt`, then run notebooks in order `01_ingest_corpus` → `02_build_index` → `03_gold_set` → `04_automl` → `05_online_learning`. Outputs regenerate in `results/` (run cards, plots) and `data/` (CSV, gold set).